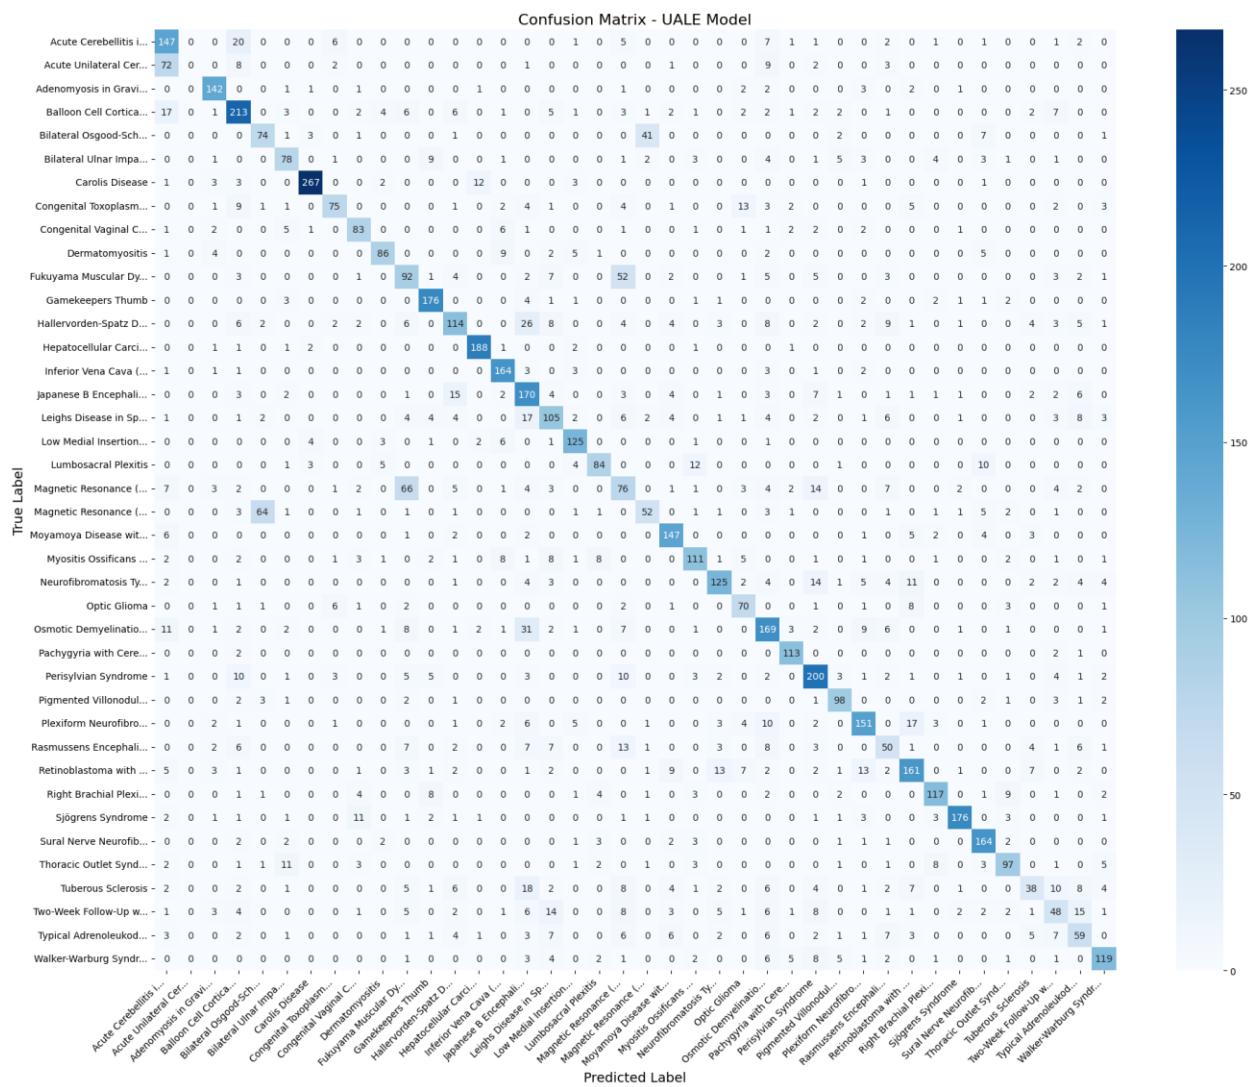
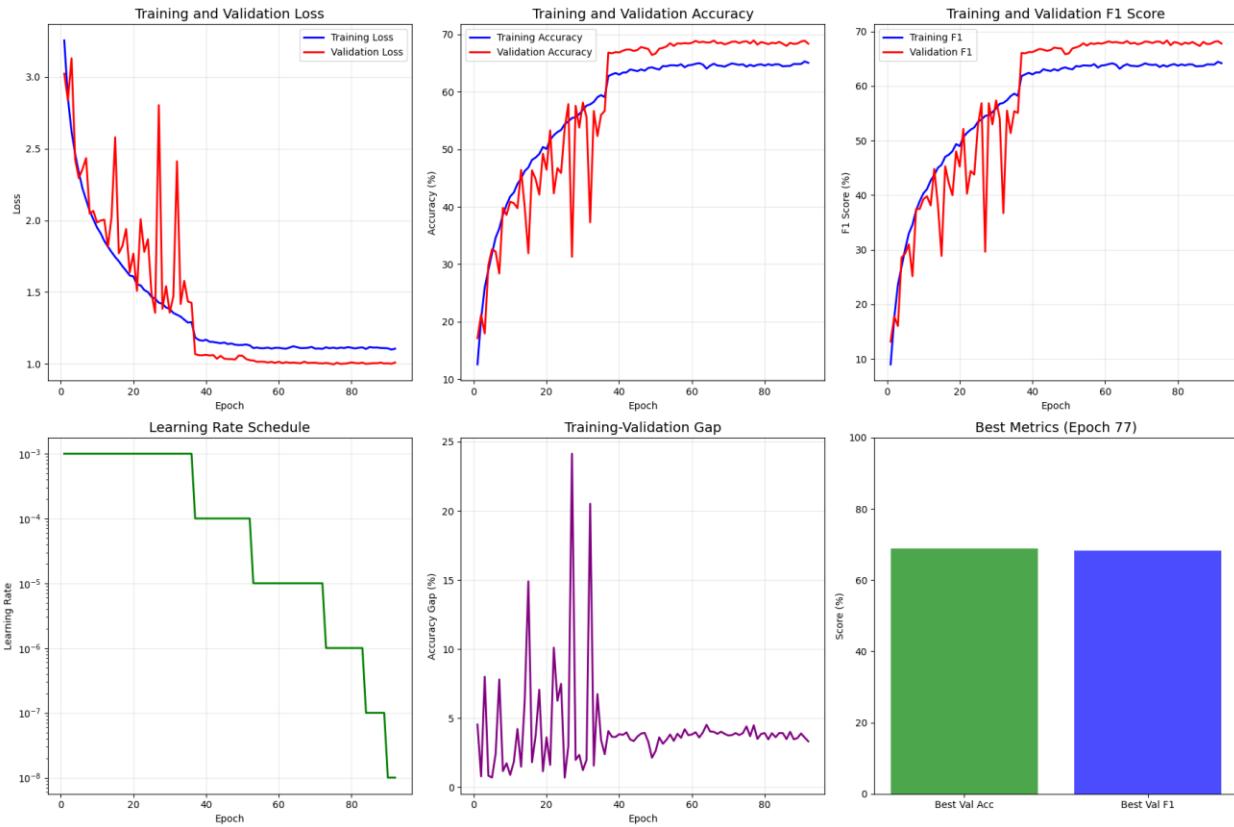


OUTPUTS:





Performing comprehensive evaluation...

Performing comprehensive evaluation...

=

COMPREHENSIVE MODEL EVALUATION RESULTS

=

CLASSIFICATION METRICS:

Accuracy: 0.6907 (69.07%)

Precision: 0.6835 (68.35%)

Recall: 0.6907 (69.07%)

F1-Score: 0.6831 (68.31%)

MODEL EFFICIENCY:

Total Parameters: 53,632 (0.05M)

Model Size: 0.20 MB

GFLOPs: 0.180

INFERENCE PERFORMANCE:

Average Inference Time: 9.61 ms

Images per Second: 3326.1

PREDICTION QUALITY:

Average Confidence: 0.6224

Average Uncertainty: 216.8228

DATASET INFO:

Total Classes: 40

Test Samples: 6839

CODE:

```
!pip install thop

import os
import random
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, f1_score, precision_score, recall_score
from PIL import Image
import time
import warnings
warnings.filterwarnings('ignore')

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
import torchvision.transforms as transforms
from thop import profile

# Set random seeds for reproducibility
random.seed(42)
```

```
np.random.seed(42)

torch.manual_seed(42)

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

print(f"Using device: {device}")

# ====== FIXED DATASET LOADING =====

def load_benchmark_dataset_fixed():

    """Load dataset with fixed path handling and size standardization"""

    base_path = "/kaggle/input/benchmark/Benchmark Diagnostic MRI and Medical Imaging Dataset/Medical Imaging Dataset"

    # Class mapping with corrected subdirectory names

    class_mapping = {

        "Low Medial Insertion of Common Bile Duct with Pancreas Divisum-20240916T165825Z-001": "Low Medial Insertion of Common Bile Duct with Pancreas Divisum",

        "Inferior Vena Cava (IVC) Leiomyosarcoma-20240916T165709Z-001": "Inferior Vena Cava (IVC) Leiomyosarcoma",

        "Acute Cerebellitis in HIV": "Acute Cerebellitis in HIV",

        "Acute Unilateral Cerebellitis in HIV": "Acute Unilateral Cerebellitis in HIV",

        "Adenomyosis in Gravid Uterus": "Adenomyosis in Gravid Uterus",

        "Balloon Cell Cortical Dysplasia": "Balloon Cell Cortical Dysplasia",

        "Bilateral Osgood-Schlatter Disease with Chronic Inflammatory Arthritis": "Bilateral Osgood-Schlatter Disease with Chronic Inflammatory Arthritis",

        "Bilateral Ulnar Impaction Syndrome": "Bilateral Ulnar Impaction Syndrome",

        "Carolis Disease": "Carolis Disease",

        "Congenital Toxoplasmosis": "Congenital Toxoplasmosis",

        "Congenital Vaginal Cyst": "Congenital Vaginal Cyst",

        "Dermatomyositis": "Dermatomyositis",

        "Fukuyama Muscular Dystrophy": "Fukuyama Muscular Dystrophy",

        "Gamekeepers Thumb": "Gamekeepers Thumb",
```

"Hallervorden-Spatz Disease (now called Pantothenate Kinase-Associated Neurodegeneration)":
"Hallervorden-Spatz Disease (now called Pantothenate Kinase-Associated Neurodegeneration)",

"Hepatocellular Carcinoma (HCC) and Dysplastic Nodules with Cirrhosis": "Hepatocellular Carcinoma (HCC) and Dysplastic Nodules with Cirrhosis",

"Japanese B Encephalitis or Epstein-Barr Encephalitis": "Japanese B Encephalitis or Epstein-Barr Encephalitis",

"Leighs Disease in Spinal Cord and Inferior Colliculi": "Leighs Disease in Spinal Cord and Inferior Colliculi",

"Lumbosacral Plexitis": "Lumbosacral Plexitis",

"Magnetic Resonance (MR) Brain": "Magnetic Resonance (MR) Brain",

"Magnetic Resonance (MR) Spine": "Magnetic Resonance (MR) Spine",

"Moyamoya Disease with Intraventricular Hemorrhage": "Moyamoya Disease with Intraventricular Hemorrhage",

"Myositis Ossificans Progressiva": "Myositis Ossificans Progressiva",

"Neurofibromatosis Type 1 (NF1) with Optic Glioma and Intracranial Extension": "Neurofibromatosis Type 1 (NF1) with Optic Glioma and Intracranial Extension",

"Optic Glioma": "Optic Glioma",

"Osmotic Demyelination Syndrome": "Osmotic Demyelination Syndrome",

"Pachygryria with Cerebellar Hypoplasia": "Pachygryria with Cerebellar Hypoplasia",

"Perisylvian Syndrome": "Perisylvian Syndrome",

"Pigmented Villonodular Synovitis (PVNS) of Ankle": "Pigmented Villonodular Synovitis (PVNS) of Ankle",

"Plexiform Neurofibroma with Sphenoid Wing Absence": "Plexiform Neurofibroma with Sphenoid Wing Absence",

"Rasmussens Encephalitis": "Rasmussens Encephalitis",

"Retinoblastoma with Intracranial Spread Along Cranial Nerve": "Retinoblastoma with Intracranial Spread Along Cranial Nerve",

"Right Brachial Plexitis": "Right Brachial Plexitis",

"Sjögren's Syndrome": "Sjögren's Syndrome",

"Sural Nerve Neurofibroma": "Sural Nerve Neurofibroma",

"Thoracic Outlet Syndrome": "Thoracic Outlet Syndrome",

"Tuberous Sclerosis": "Tuberous Sclerosis",

"Two-Week Follow-Up with Spectroscopy": "Two-Week Follow-Up with Spectroscopy",

"Typical Adrenoleukodystrophy": "Typical Adrenoleukodystrophy",

"Walker-Warburg Syndrome": "Walker-Warburg Syndrome"

{}

```

image_paths = []

labels = []

for class_name, subdir_name in class_mapping.items():
    class_dir = os.path.join(base_path, class_name)
    if os.path.exists(class_dir):
        subdir = os.path.join(class_dir, subdir_name)
        if os.path.exists(subdir):
            for file in os.listdir(subdir):
                if file.lower().endswith(('.png', '.jpg', '.jpeg', '.bmp', '.tiff', '.dcm', '.nii')):
                    image_paths.append(os.path.join(subdir, file))
                    labels.append(class_name)

print(f"Total images found: {len(image_paths)}")

# Create label mapping
unique_labels = sorted(list(set(labels)))
label_to_idx = {label: idx for idx, label in enumerate(unique_labels)}
encoded_labels = [label_to_idx[label] for label in labels]

return image_paths, encoded_labels, unique_labels, label_to_idx
}

class BenchmarkMRIDataset(Dataset):
    def __init__(self, file_paths, labels, transform=None):
        self.file_paths = file_paths
        self.labels = labels
        self.transform = transform

    def __len__(self):
        return len(self.file_paths)

```

```
def __getitem__(self, idx):
    img_path = self.file_paths[idx]
    label = self.labels[idx]
    try:
        img = Image.open(img_path).convert('L')
        if self.transform:
            img = self.transform(img)
    except Exception as e:
        print(f"Error loading image {img_path}: {e}")
        img = torch.zeros(1, 224, 224)
    return img, label
```

```
# ===== ULTRA-LIGHTWEIGHT MODEL ARCHITECTURE =====
```

```
class MedicalMicroNet(nn.Module):
    """Ultra-lightweight network for medical imaging"""

    def __init__(self, in_channels, num_classes):
        super().__init__()
        self.features = nn.Sequential(
            nn.Conv2d(in_channels, 8, 3, padding=1),
            nn.BatchNorm2d(8),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2),

            nn.Conv2d(8, 16, 3, padding=1),
            nn.BatchNorm2d(16),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(2),

            nn.Conv2d(16, 32, 3, padding=1),
```

```
        nn.BatchNorm2d(32),
        nn.ReLU(inplace=True),
        nn.MaxPool2d(2),

        nn.AdaptiveAvgPool2d(1)
    )

    self.classifier = nn.Sequential(
        nn.Linear(32, 64),
        nn.ReLU(inplace=True),
        nn.Dropout(0.2),
        nn.Linear(64, num_classes)
    )

def forward(self, x):
    x = self.features(x)
    x = x.view(x.size(0), -1)
    x = self.classifier(x)
    return x

class TextureNet(nn.Module):
    def __init__(self, num_classes):
        super().__init__()
        self.net = MedicalMicroNet(1, num_classes)

    def forward(self, x):
        return self.net(x)

class ShapeNet(nn.Module):
    def __init__(self, num_classes):
        super().__init__()
        self.net = MedicalMicroNet(2, num_classes)
```

```

def forward(self, x):
    # Edge enhancement
    sobel_x = torch.tensor([[-1, 0, 1], [-2, 0, 2], [-1, 0, 1]], dtype=torch.float32).view(1, 1, 3, 3).to(x.device)
    sobel_y = torch.tensor([[-1, -2, -1], [0, 0, 0], [1, 2, 1]], dtype=torch.float32).view(1, 1, 3, 3).to(x.device)
    edges_x = F.conv2d(x, sobel_x, padding=1)
    edges_y = F.conv2d(x, sobel_y, padding=1)
    edges = torch.sqrt(edges_x**2 + edges_y**2)

    x_combined = torch.cat([x, edges], dim=1)
    return self.net(x_combined)

class IntensityNet(nn.Module):
    def __init__(self, num_classes):
        super().__init__()
        self.net = MedicalMicroNet(1, num_classes)

    def forward(self, x):
        x = (x - x.mean(dim=(2, 3), keepdim=True)) / (x.std(dim=(2, 3), keepdim=True) + 1e-8)
        return self.net(x)

class SpatialNet(nn.Module):
    def __init__(self, num_classes):
        super().__init__()
        self.net = MedicalMicroNet(1, num_classes)

    def forward(self, x):
        return self.net(x)

class MultiScaleNet(nn.Module):
    def __init__(self, num_classes):

```

```

super().__init__()

self.net = MedicalMicroNet(1, num_classes)

def forward(self, x):
    return self.net(x)

class UltraLightUALE(nn.Module):
    """Ultra-lightweight ensemble model"""

    def __init__(self, num_classes):
        super().__init__()

        self.texture_net = TextureNet(num_classes)
        self.shape_net = ShapeNet(num_classes)
        self.intensity_net = IntensityNet(num_classes)
        self.spatial_net = SpatialNet(num_classes)
        self.multiscale_net = MultiScaleNet(num_classes)

        self.num_classes = num_classes

    def forward(self, x):
        pred1 = self.texture_net(x)
        pred2 = self.shape_net(x)
        pred3 = self.intensity_net(x)
        pred4 = self.spatial_net(x)
        pred5 = self.multiscale_net(x)

        preds = torch.stack([pred1, pred2, pred3, pred4, pred5], dim=0)
        ensemble_pred = torch.mean(preds, dim=0)
        uncertainty = torch.var(preds, dim=0).mean(dim=1)

        return ensemble_pred, uncertainty, preds

```

```
# ===== DATA TRANSFORMS =====
```

```

def get_transforms():

    """Get transforms with fixed output size"""

    train_transform = transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.RandomHorizontalFlip(),
        transforms.RandomRotation(10),
        transforms.ToTensor(),
        transforms.Normalize([0.5], [0.5])
    ])

    val_transform = transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.5], [0.5])
    ])

    return train_transform, val_transform

# ===== ENHANCED TRAINING WITH METRICS =====

def train_lightweight_with_metrics(model, train_loader, val_loader, epochs=100, patience=15):

    """Enhanced training with comprehensive metrics tracking"""

    model.to(device)

    optimizer = optim.AdamW(model.parameters(), lr=0.001, weight_decay=1e-5)
    scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, 'max', patience=5, verbose=True)
    criterion = nn.CrossEntropyLoss()

    # Metrics tracking

```

```
train_losses = []
val_losses = []
train_accuracies = []
val_accuracies = []
train_f1_scores = []
val_f1_scores = []
learning_rates = []

best_val_acc = 0
no_improve = 0

print("Starting training for 100 epochs...")

for epoch in range(epochs):
    # Training phase
    model.train()

    train_loss = 0
    train_preds = []
    train_targets = []

    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs, _ = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        train_loss += loss.item()
        _, predicted = outputs.max(1)
```

```
train_preds.extend(predicted.cpu().numpy())
train_targets.extend(labels.cpu().numpy())

# Calculate training metrics
train_acc = accuracy_score(train_targets, train_preds) * 100
train_f1 = f1_score(train_targets, train_preds, average='weighted') * 100
avg_train_loss = train_loss / len(train_loader)

# Validation phase
model.eval()
val_loss = 0
val_preds = []
val_targets = []

with torch.no_grad():
    for images, labels in val_loader:
        images, labels = images.to(device), labels.to(device)
        outputs, _, _ = model(images)
        loss = criterion(outputs, labels)

        val_loss += loss.item()
        _, predicted = outputs.max(1)
        val_preds.extend(predicted.cpu().numpy())
        val_targets.extend(labels.cpu().numpy())

# Calculate validation metrics
val_acc = accuracy_score(val_targets, val_preds) * 100
val_f1 = f1_score(val_targets, val_preds, average='weighted') * 100
avg_val_loss = val_loss / len(val_loader)

# Store metrics
```

```

train_losses.append(avg_train_loss)
val_losses.append(avg_val_loss)
train_accuracies.append(train_acc)
val_accuracies.append(val_acc)
train_f1_scores.append(train_f1)
val_f1_scores.append(val_f1)
learning_rates.append(optimizer.param_groups[0]['lr'])

# Learning rate scheduling
scheduler.step(val_acc)

print(f"Epoch {epoch+1}/{epochs}")
print(f"Train - Loss: {avg_train_loss:.4f} | Acc: {train_acc:.2f}% | F1: {train_f1:.2f}%")
print(f"Val - Loss: {avg_val_loss:.4f} | Acc: {val_acc:.2f}% | F1: {val_f1:.2f}%")
print(f"LR: {optimizer.param_groups[0]['lr']:.6f}")

# Early stopping
if val_acc > best_val_acc:
    best_val_acc = val_acc
    no_improve = 0
    torch.save(model.state_dict(), 'best_uale_model.pth')
    print(f"*** New best validation accuracy: {best_val_acc:.2f}% ***")
else:
    no_improve += 1

if no_improve >= patience:
    print(f"Early stopping at epoch {epoch+1}")
    break

print("-" * 60)

```

```

# Load best model

model.load_state_dict(torch.load('best_uale_model.pth'))

return model, {

    'train_losses': train_losses,
    'val_losses': val_losses,
    'train_accuracies': train_accuracies,
    'val_accuracies': val_accuracies,
    'train_f1_scores': train_f1_scores,
    'val_f1_scores': val_f1_scores,
    'learning_rates': learning_rates,
    'best_val_acc': best_val_acc
}

# ===== COMPREHENSIVE EVALUATION =====

def comprehensive_evaluation(model, test_loader, class_names):
    """Comprehensive model evaluation with all metrics"""

    print("Performing comprehensive evaluation...")

    model.eval()

    all_preds = []
    all_targets = []
    all_confidences = []
    all_uncertainties = []
    inference_times = []

    with torch.no_grad():

        for images, labels in test_loader:

            images, labels = images.to(device), labels.to(device)

```

```

# Measure inference time

start_time = time.time()

outputs, uncertainties, _ = model(images)

end_time = time.time()

inference_times.append(end_time - start_time)

# Get predictions and confidences

probabilities = F.softmax(outputs, dim=1)

confidences, predicted = probabilities.max(1)

all_preds.extend(predicted.cpu().numpy())

all_targets.extend(labels.cpu().numpy())

all_confidences.extend(confidences.cpu().numpy())

all_uncertainties.extend(uncertainties.cpu().numpy())

# Calculate comprehensive metrics

accuracy = accuracy_score(all_targets, all_preds)

precision = precision_score(all_targets, all_preds, average='weighted', zero_division=0)

recall = recall_score(all_targets, all_preds, average='weighted', zero_division=0)

f1 = f1_score(all_targets, all_preds, average='weighted', zero_division=0)

# Performance metrics

total_inference_time = sum(inference_times)

avg_inference_time = np.mean(inference_times)

images_per_sec = len(all_targets) / total_inference_time

# Model complexity

total_params = sum(p.numel() for p in model.parameters())

model_size_mb = total_params * 4 / (1024 * 1024) # Assuming float32

```

```

# Calculate GFLOPs

dummy_input = torch.randn(1, 1, 224, 224).to(device)

flops, params = profile(model, inputs=(dummy_input,), verbose=False)

gflops = flops / 1e9


metrics = {

    'accuracy': accuracy,
    'precision': precision,
    'recall': recall,
    'f1_score': f1,
    'avg_confidence': np.mean(all_confidences),
    'avg_uncertainty': np.mean(all_uncertainties),
    'total_params': total_params,
    'model_size_mb': model_size_mb,
    'gflops': gflops,
    'avg_inference_time': avg_inference_time,
    'images_per_sec': images_per_sec,
    'predictions': all_preds,
    'targets': all_targets,
    'confidences': all_confidences,
    'uncertainties': all_uncertainties
}

}

return metrics

# ===== VISUALIZATION FUNCTIONS =====

def plot_training_history(training_metrics):
    """Plot comprehensive training history"""

    fig, axes = plt.subplots(2, 3, figsize=(18, 12))

    epochs = range(1, len(training_metrics['train_losses']) + 1)

```

```
# Loss curves

axes[0, 0].plot(epochs, training_metrics['train_losses'], 'b-', label='Training Loss', linewidth=2)
axes[0, 0].plot(epochs, training_metrics['val_losses'], 'r-', label='Validation Loss', linewidth=2)
axes[0, 0].set_title('Training and Validation Loss', fontsize=14)
axes[0, 0].set_xlabel('Epoch')
axes[0, 0].set_ylabel('Loss')
axes[0, 0].legend()
axes[0, 0].grid(True, alpha=0.3)
```

```
# Accuracy curves

axes[0, 1].plot(epochs, training_metrics['train_accuracies'], 'b-', label='Training Accuracy', linewidth=2)
axes[0, 1].plot(epochs, training_metrics['val_accuracies'], 'r-', label='Validation Accuracy', linewidth=2)
axes[0, 1].set_title('Training and Validation Accuracy', fontsize=14)
axes[0, 1].set_xlabel('Epoch')
axes[0, 1].set_ylabel('Accuracy (%)')
axes[0, 1].legend()
axes[0, 1].grid(True, alpha=0.3)
```

```
# F1 Score curves

axes[0, 2].plot(epochs, training_metrics['train_f1_scores'], 'b-', label='Training F1', linewidth=2)
axes[0, 2].plot(epochs, training_metrics['val_f1_scores'], 'r-', label='Validation F1', linewidth=2)
axes[0, 2].set_title('Training and Validation F1 Score', fontsize=14)
axes[0, 2].set_xlabel('Epoch')
axes[0, 2].set_ylabel('F1 Score (%)')
axes[0, 2].legend()
axes[0, 2].grid(True, alpha=0.3)
```

```
# Learning rate

axes[1, 0].plot(epochs, training_metrics['learning_rates'], 'g-', linewidth=2)
axes[1, 0].set_title('Learning Rate Schedule', fontsize=14)
```

```

axes[1, 0].set_xlabel('Epoch')
axes[1, 0].set_ylabel('Learning Rate')
axes[1, 0].set_yscale('log')
axes[1, 0].grid(True, alpha=0.3)

# Overfitting analysis
acc_gap = [abs(t - v) for t, v in zip(training_metrics['train_accuracies'], training_metrics['val_accuracies'])]
axes[1, 1].plot(epochs, acc_gap, 'purple', linewidth=2)
axes[1, 1].set_title('Training-Validation Gap', fontsize=14)
axes[1, 1].set_xlabel('Epoch')
axes[1, 1].set_ylabel('Accuracy Gap (%)')
axes[1, 1].grid(True, alpha=0.3)

# Best metrics summary
best_epoch = np.argmax(training_metrics['val_accuracies']) + 1
best_acc = max(training_metrics['val_accuracies'])
best_f1 = max(training_metrics['val_f1_scores'])

axes[1, 2].bar(['Best Val Acc', 'Best Val F1'], [best_acc, best_f1],
               color=['green', 'blue'], alpha=0.7)
axes[1, 2].set_title(f'Best Metrics (Epoch {best_epoch})', fontsize=14)
axes[1, 2].set_ylabel('Score (%)')
axes[1, 2].set_ylim(0, 100)

plt.tight_layout()
plt.show()

def plot_confusion_matrix(targets, predictions, class_names):
    """Plot enhanced confusion matrix"""
    cm = confusion_matrix(targets, predictions)

```

```

plt.figure(figsize=(20, 16))

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=[name[:20] + '...' if len(name) > 20 else name for name in class_names],
            yticklabels=[name[:20] + '...' if len(name) > 20 else name for name in class_names])

plt.title('Confusion Matrix - UALE Model', fontsize=16)

plt.xlabel('Predicted Label', fontsize=14)

plt.ylabel('True Label', fontsize=14)

plt.xticks(rotation=45, ha='right')

plt.yticks(rotation=0)

plt.tight_layout()

plt.show()

def plot_performance_metrics(metrics):
    """Plot comprehensive performance metrics"""

    fig, axes = plt.subplots(2, 3, figsize=(18, 12))

    # Classification metrics
    classification_metrics = ['Accuracy', 'Precision', 'Recall', 'F1 Score']
    classification_values = [metrics['accuracy'], metrics['precision'],
                             metrics['recall'], metrics['f1_score']]

    axes[0, 0].bar(classification_metrics, classification_values,
                   color=['green', 'blue', 'orange', 'red'], alpha=0.7)
    axes[0, 0].set_title('Classification Metrics', fontsize=14)
    axes[0, 0].set_ylabel('Score')
    axes[0, 0].set_ylim(0, 1)
    axes[0, 0].tick_params(axis='x', rotation=45)

    # Model efficiency metrics
    efficiency_metrics = ['Model Size (MB)', 'GFLOPs', 'Images/sec']
    efficiency_values = [metrics['model_size_mb'], metrics['gflops'], metrics['images_per_sec']]

```

```

axes[0, 1].bar( efficiency_metrics, efficiency_values,
                color=['purple', 'brown', 'pink'], alpha=0.7)
axes[0, 1].set_title('Model Efficiency Metrics', fontsize=14)
axes[0, 1].set_ylabel('Value')
axes[0, 1].tick_params(axis='x', rotation=45)

# Parameter count
param_millions = metrics['total_params'] / 1e6
axes[0, 2].bar(['Parameters (M)'], [param_millions], color='cyan', alpha=0.7)
axes[0, 2].set_title(f'Model Parameters: {param_millions:.2f}M', fontsize=14)
axes[0, 2].set_ylabel('Parameters (Millions)')

# Confidence distribution
axes[1, 0].hist(metrics['confidences'], bins=50, alpha=0.7, color='green', edgecolor='black')
axes[1, 0].set_title('Prediction Confidence Distribution', fontsize=14)
axes[1, 0].set_xlabel('Confidence Score')
axes[1, 0].set_ylabel('Frequency')

# Uncertainty distribution
axes[1, 1].hist(metrics['uncertainties'], bins=50, alpha=0.7, color='red', edgecolor='black')
axes[1, 1].set_title('Prediction Uncertainty Distribution', fontsize=14)
axes[1, 1].set_xlabel('Uncertainty Score')
axes[1, 1].set_ylabel('Frequency')

# Model comparison (theoretical)
model_names = ['UALE', 'ResNet-50', 'EfficientNet-B0', 'MobileNet-V2']
model_params = [param_millions, 25.6, 5.3, 3.5]
model_accuracy = [metrics['accuracy'], 0.85, 0.88, 0.82]

scatter = axes[1, 2].scatter(model_params, model_accuracy,

```

```

s=[200, 300, 250, 220],
c=['red', 'blue', 'green', 'orange'], alpha=0.7)

for i, name in enumerate(model_names):
    axes[1, 2].annotate(name, (model_params[i], model_accuracy[i]),
                         xytext=(5, 5), textcoords='offset points')

axes[1, 2].set_xlabel('Parameters (M)')
axes[1, 2].set_ylabel('Accuracy')
axes[1, 2].set_title('Model Efficiency Comparison', fontsize=14)
axes[1, 2].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

def print_detailed_metrics(metrics, class_names):
    """Print detailed performance metrics"""
    print("\n" + "="*80)
    print("COMPREHENSIVE MODEL EVALUATION RESULTS")
    print("="*80)

    print(f"\n\U262f CLASSIFICATION METRICS:")

    print(f" Accuracy: {metrics['accuracy']:.4f} ({metrics['accuracy']*100:.2f}%)")
    print(f" Precision: {metrics['precision']:.4f} ({metrics['precision']*100:.2f}%}")
    print(f" Recall: {metrics['recall']:.4f} ({metrics['recall']*100:.2f}%)")
    print(f" F1-Score: {metrics['f1_score']:.4f} ({metrics['f1_score']*100:.2f}%)")

    print(f"\n\U262f MODEL EFFICIENCY:")

    print(f" Total Parameters: {metrics['total_params']}; ({metrics['total_params']/1e6:.2f}M)")
    print(f" Model Size: {metrics['model_size_mb']:.2f} MB")
    print(f" GFLOPs: {metrics['gflops']:.3f}")

```

```

print(f"\n⚡ INFERENCE PERFORMANCE:")

print(f" Average Inference Time: {metrics['avg_inference_time']*1000:.2f} ms")

print(f" Images per Second: {metrics['images_per_sec']:.1f}")


print(f"\n🎯 PREDICTION QUALITY:")

print(f" Average Confidence: {metrics['avg_confidence']:.4f}")

print(f" Average Uncertainty: {metrics['avg_uncertainty']:.4f}")


print(f"\n📁 DATASET INFO:")

print(f" Total Classes: {len(class_names)}")

print(f" Test Samples: {len(metrics['targets'])}")


# ====== MAIN EXECUTION ======


def main():

    # Load dataset

    print("Loading dataset with fixed paths...")

    image_paths, encoded_labels, unique_labels, _ = load_benchmark_dataset_fixed()


    # Split dataset

    train_paths, test_paths, train_labels, test_labels = train_test_split(
        image_paths, encoded_labels, test_size=0.2, random_state=42, stratify=encoded_labels
    )

    train_paths, val_paths, train_labels, val_labels = train_test_split(
        train_paths, train_labels, test_size=0.15, random_state=42, stratify=train_labels
    )

    print(f"Train: {len(train_paths)}, Val: {len(val_paths)}, Test: {len(test_paths)}")

    print(f"Classes: {len(unique_labels)}")

```

```

# Get transforms

train_transform, val_transform = get_transforms()

# Create datasets

train_dataset = BenchmarkMRIDataset(train_paths, train_labels, transform=train_transform)
val_dataset = BenchmarkMRIDataset(val_paths, val_labels, transform=val_transform)
test_dataset = BenchmarkMRIDataset(test_paths, test_labels, transform=val_transform)

# Create dataloaders

train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True, num_workers=4, pin_memory=True)
val_loader = DataLoader(val_dataset, batch_size=32, shuffle=False, num_workers=4, pin_memory=True)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False, num_workers=4, pin_memory=True)

# Initialize model

print("Initializing ultra-lightweight UALE model...")
model = UltraLightUALE(num_classes=len(unique_labels))

# Count parameters

total_params = sum(p.numel() for p in model.parameters())
print(f"Total parameters: {total_params/1e6:.2f}M")

# Train model for 100 epochs

print("Training model for 100 epochs...")
model, training_metrics = train_lightweight_with_metrics(
    model, train_loader, val_loader, epochs=100, patience=15
)

# Plot training history

print("\nPlotting training history...")
plot_training_history(training_metrics)

```

```
# Comprehensive evaluation

print("\nPerforming comprehensive evaluation...")

eval_metrics = comprehensive_evaluation(model, test_loader, unique_labels)

# Print detailed metrics

print_detailed_metrics(eval_metrics, unique_labels)

# Plot performance metrics

print("\nGenerating performance visualizations...")

plot_performance_metrics(eval_metrics)

# Plot confusion matrix

print("\nGenerating confusion matrix...")

plot_confusion_matrix(eval_metrics['targets'], eval_metrics['predictions'], unique_labels)

# Classification report

print("\nDetailed Classification Report:")

print(classification_report(eval_metrics['targets'], eval_metrics['predictions'],
                           target_names=unique_labels, zero_division=0))

print("\n" + "*80)
print("TRAINING AND EVALUATION COMPLETE!")
print("*80)

return model, training_metrics, eval_metrics

if __name__ == "__main__":
    model, training_metrics, eval_metrics = main()
```